
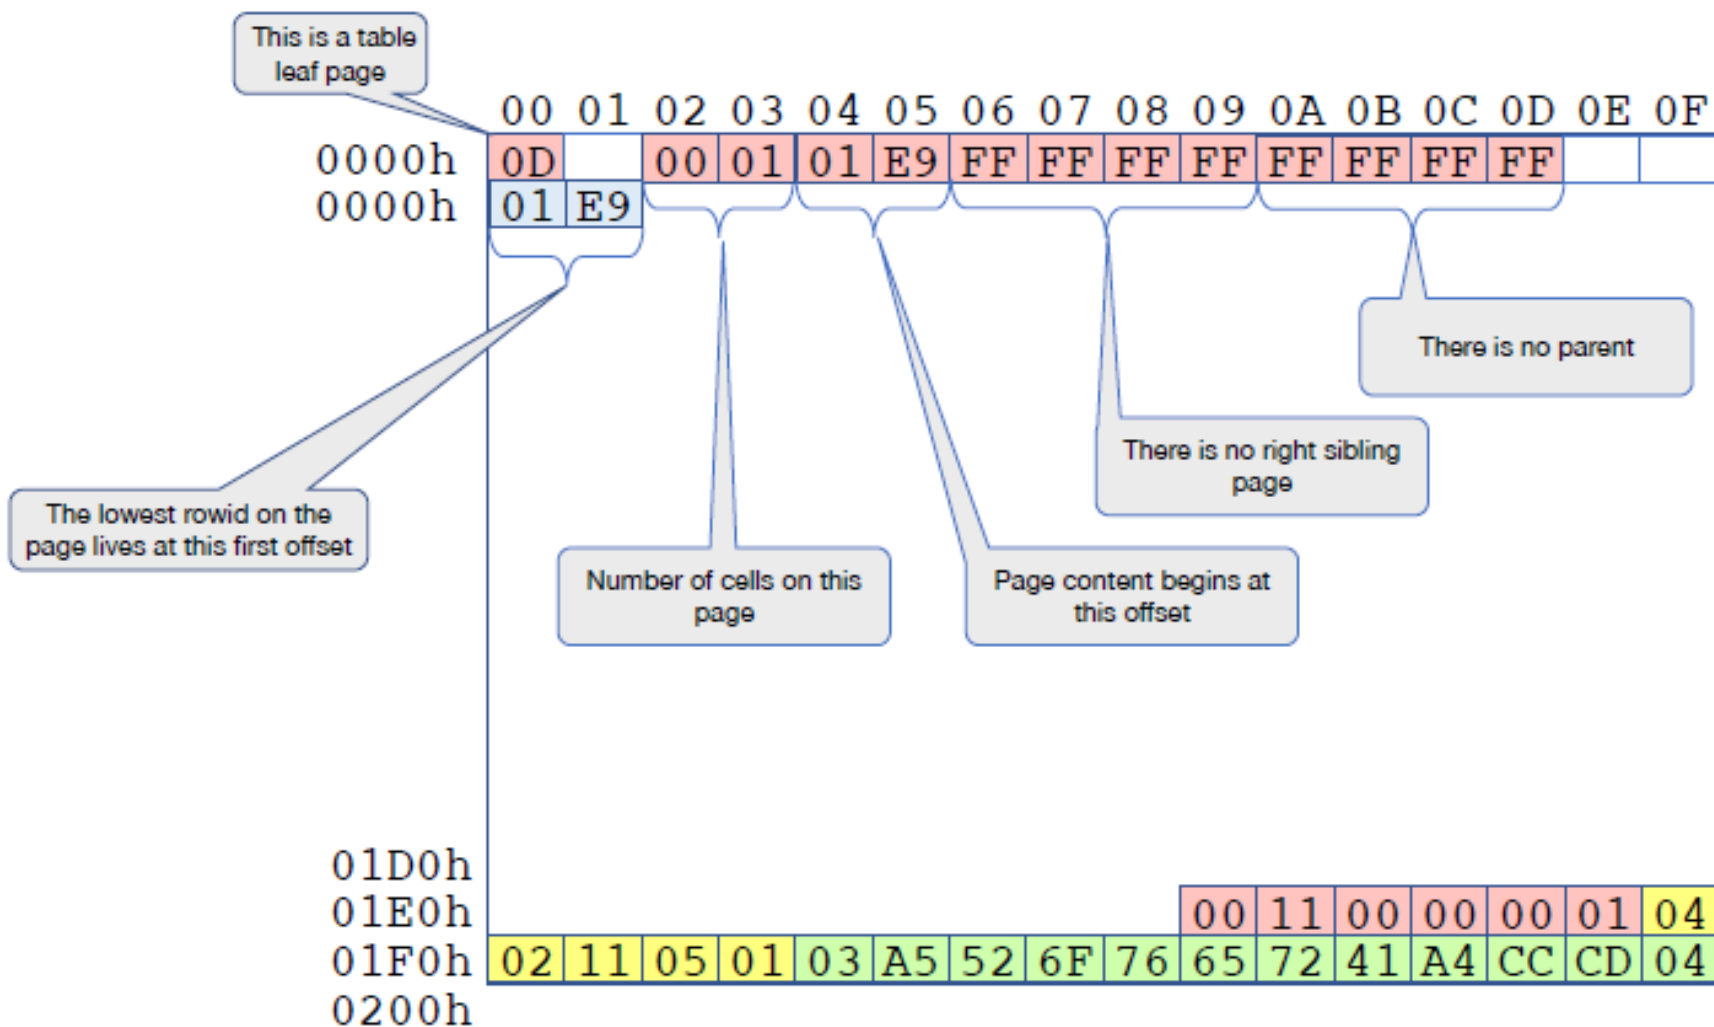
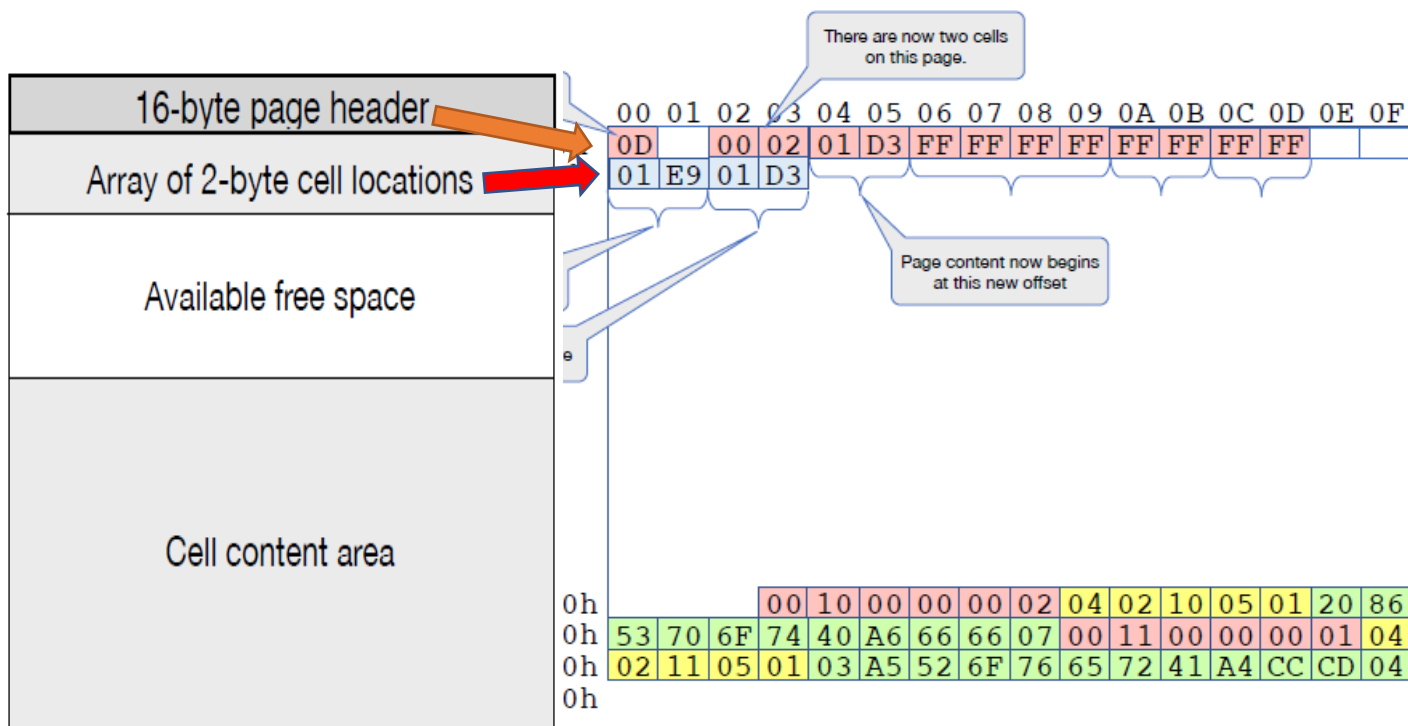


Database design guide & Storage mapping

Page Format

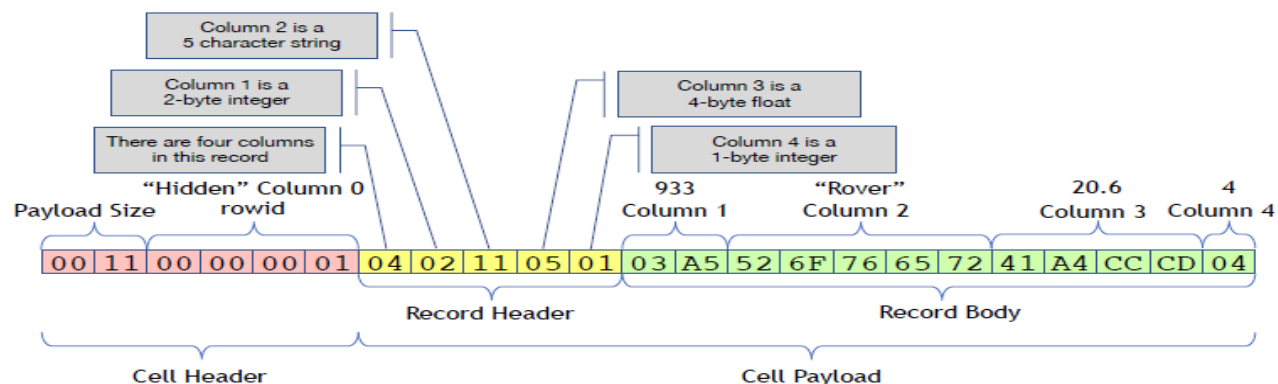




3.1. Page Headers

| Page Offset | Element Size | Description |
|-------------|--------------------------|--|
| 0x00 | 1 | The one-byte flag at page offset 0 indicates the b-tree page type. <ul style="list-style-type: none"> A value of 2 (0x02) means the page is an index b-tree interior page. A value of 5 (0x05) means the page is an table b-tree interior page. A value of 10 (0x0a) means the page is an index b-tree leaf page. A value of 13 (0x0d) means the page is a table b-tree leaf page. Any other value for the b-tree page type is an error. |
| 0x01 | 1 | Unused |
| 0x02 | 2 | The two-byte integer at offset 2 designates the number of cells on the page. |
| 0x04 | 2 | The two-byte integer at offset 4 designates the page offset for the start of the cell content area. A zero value for this integer is interpreted as 65536. |
| 0x06 | 4 | The four-byte integer page pointer at offset 0x06 has a different role depending on the b-tree page type: <ul style="list-style-type: none"> Table or Index interior page - page number of rightmost child Table or Index leaf page - page number of sibling to the right (This number is 0xFFFFFFFF if no pointer exists) |
| 0x0A | 4 | The four-byte integer page pointer at offset 0x0A references the page's parent. If this is a root page, then the special value 0xFFFFFFFF is used. |
| 0x0E | 2 | Unused |
| 0x10 | 2 x <i>cells on page</i> | An array of 2-byte integers that indicate the page offset location of each data cell. The array size is 2 <i>n</i> , where <i>n</i> is the number of cells on the page. The array is maintained in key-sorted order. |

First Record Format Insert



First Record Format Insert

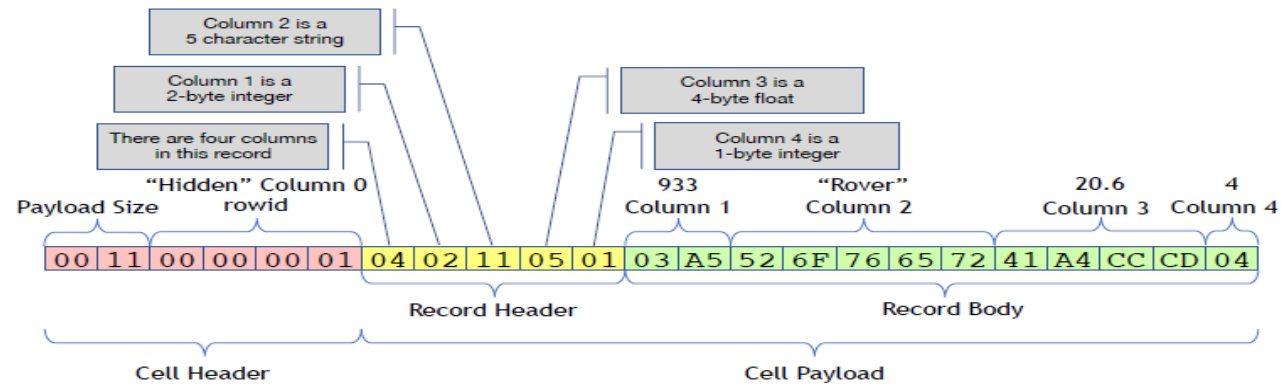
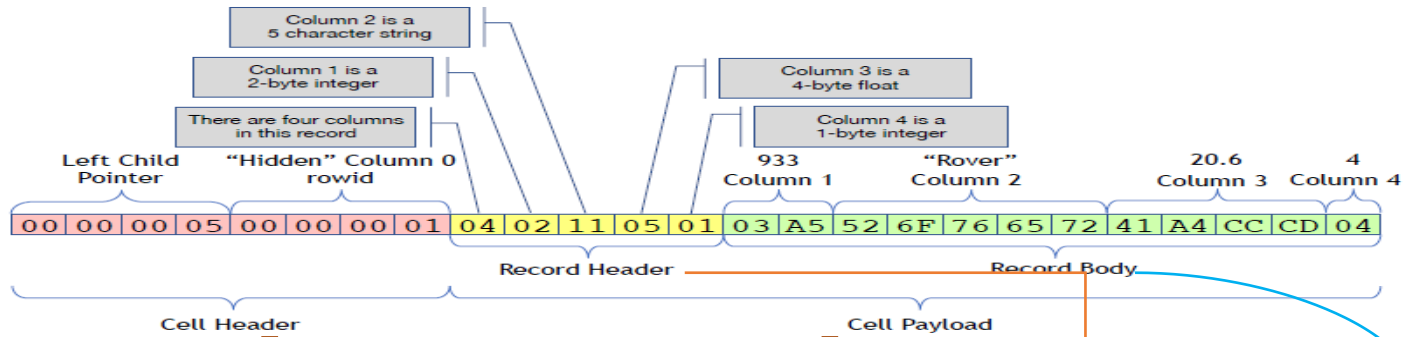


Table Interior Cell



| File Type | Page Type | Cell Header | | | Cell Body |
|-----------|-----------|------------------|-----------------------|------------|--------------|
| | | 4-byte int | 2-byte int | 4-byte int | N-byte array |
| | | Left Child Page# | Bytes of Cell Payload | Rowid | Payload |
| Table | Leaf | | ✓ | ✓ | ✓ |
| | Interior | ✓ | | ✓ | |
| Index | Leaf | | ✓ | | ✓ |
| | Interior | ✓ | ✓ | | ✓ |

Figure 2: Cell Formats

| Hex | Dec | Data Type Name | Content Size (bytes) | Description |
|----------|--------|-----------------|----------------------|--|
| 0x00 | 0 | NULL | 0 | Value is a NULL (i.e. it takes no memory in the record body) |
| 0x01 | 1 | TINYINT, BYTE | 1 | Value is a signed, 8-bit twos-complement integer. |
| 0x02 | 2 | SMALLINT, SHORT | 2 | Value is a signed, big-endian, 16-bit twos-complement integer. |
| 0x03 | 3 | INT, INTEGER | 4 | Value is a signed, big-endian, 32-bit twos-complement integer. |
| 0x04 | 4 | BIGINT, LONG | 8 | Value is a signed, big-endian, 64-bit twos-complement integer. |
| 0x05 | 5 | FLOAT | 4 | Value is a single precision, big-endian, IEEE 754-2008 32-bit floating point number. |
| 0x06 | 6 | DOUBLE, REAL | 8 | Value is a double precision, big-endian, IEEE 754-2008 64-bit floating point number. |
| 0x08 | 8 | YEAR | 1 | Value is an 8-bit twos-complement integer. Both positive and negative numbers are supported in the range -128 to 127. This indicates a year with respect to the year 2000, i.e. 1872-2127 |
| 0x09 | 9 | TIME | 4 | Value is a big-endian 32-bit twos-complement integer. Indicates time of day in milliseconds since midnight, i.e. "millis". Note that only values of [0-86,400,000) are valid. i.e. Hex [0x00000000-0x005c00) |
| 0x0A | 10 | DATETIME | 8 | A big-endian unsigned LONG integer that represents the specified number of milliseconds since the standard base time known as "the epoch". It should display as a formatted string: YYYY-MM-DD_hh:mm:ss, e.g. 2016-03-23 13:52:23. |
| 0x0B | 11 | DATE | 8 | A datetime whose time component is 00:00:00, but does not display. |
| 0x0C + n | 12 + n | TEXT | | Value is an ASCII string of length n. C-style string null terminators are not used or needed. |

Figure 4 - Data Types, Their 1-byte Code, and Implementation


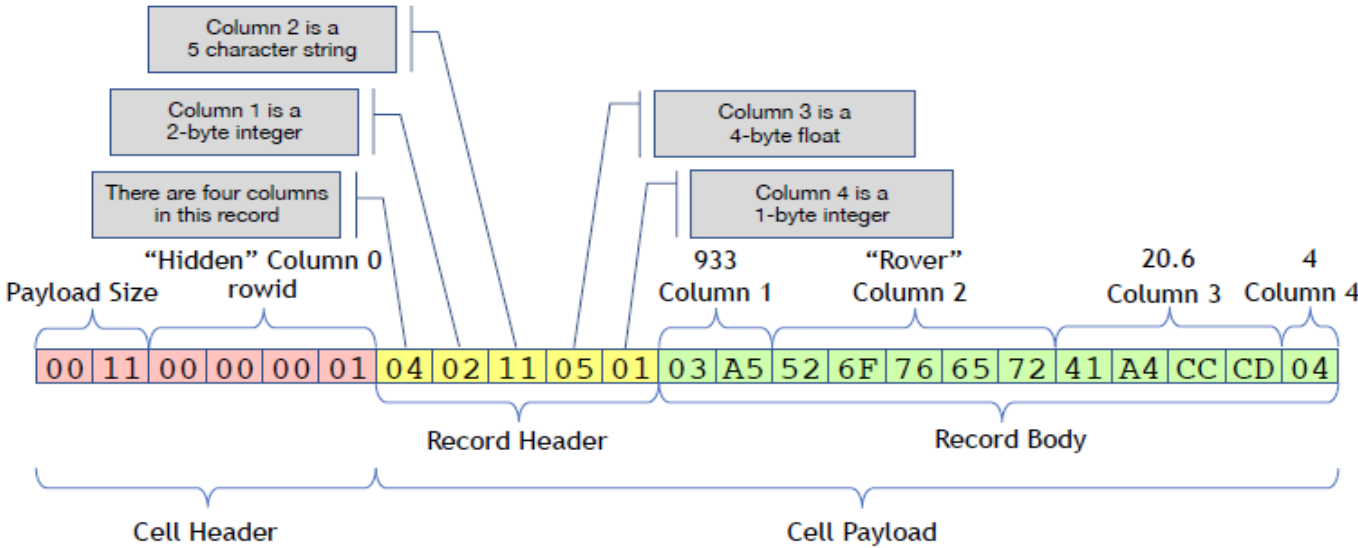


Figure 3: Record Format (i.e. Payload of Table Leaf Cell)

Example

```
CREATE TABLE Dogs (  
  Id      SHORT,  
  Name    TEXT,  
  Weight  FLOAT,  
  Age     BYTE);
```

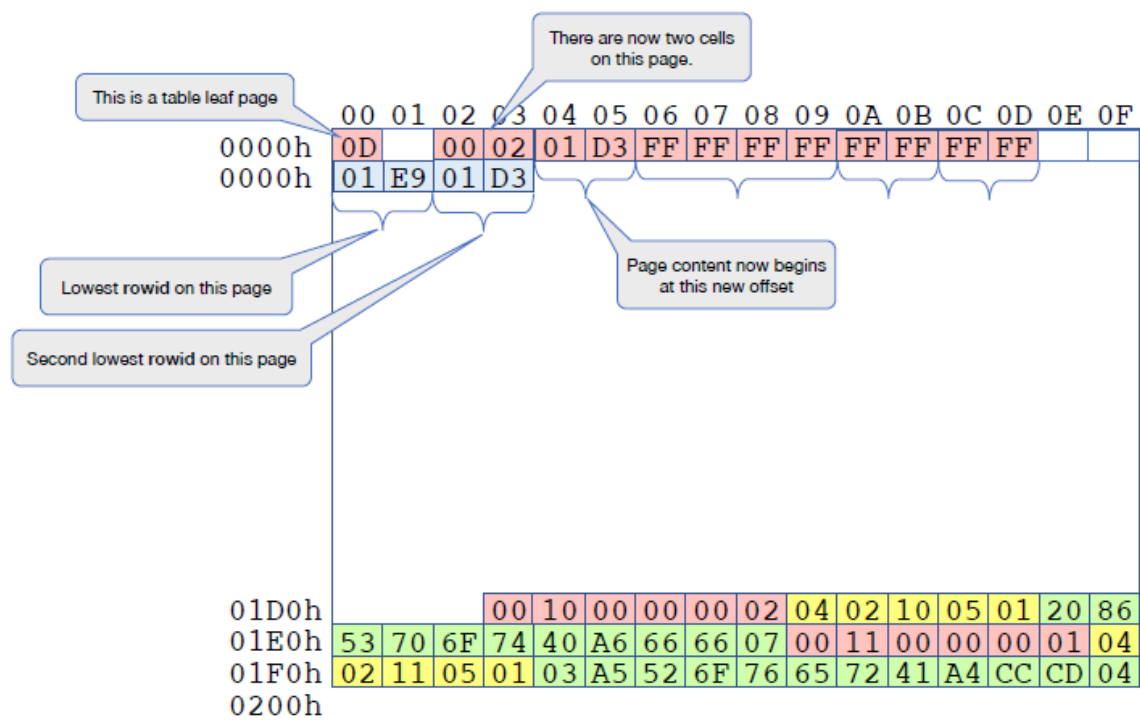
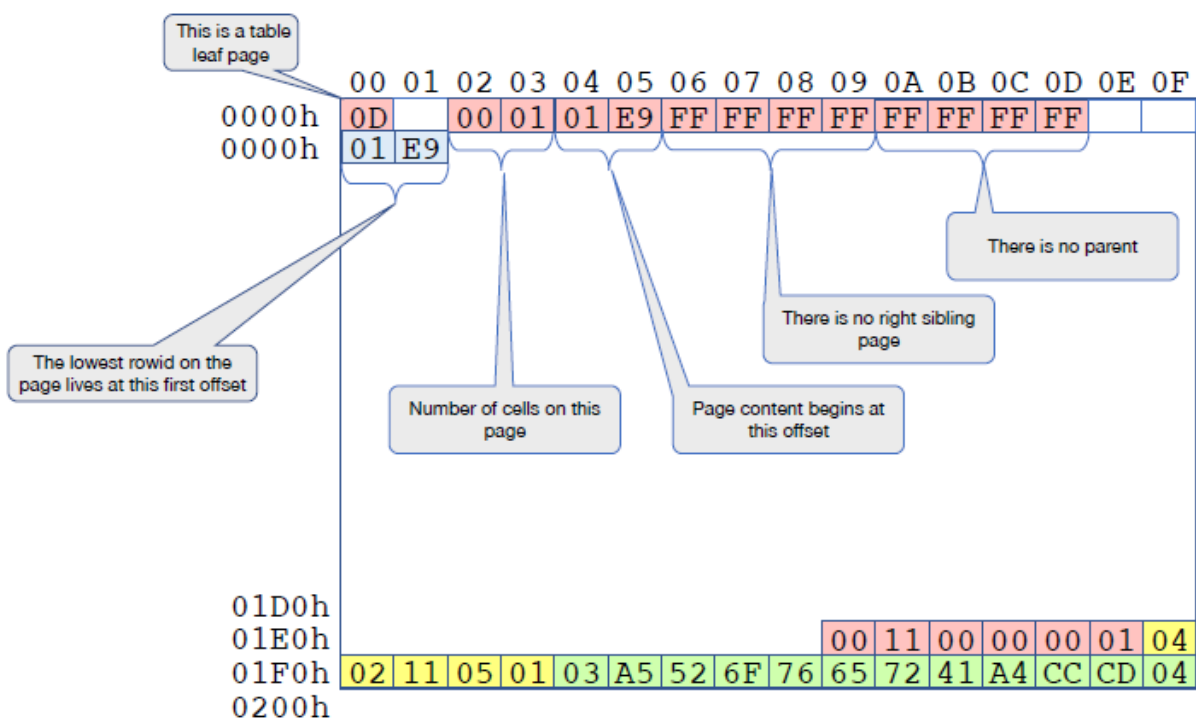
```
INSERT INTO TABLE Dogs VALUES  
(933, "Rover", 20.6, 4);
```



| Hex | Dec | Data Type Name | Content Size (bytes) | Description |
|----------|--------|-----------------|----------------------|--|
| 0x00 | 0 | NULL | 0 | Value is a NULL (i.e. it takes no memory in the record body) |
| 0x01 | 1 | TINYINT, BYTE | 1 | Value is a signed, 8-bit twos-complement integer. |
| 0x02 | 2 | SMALLINT, SHORT | 2 | Value is a signed, big-endian, 16-bit twos-complement integer. |
| 0x03 | 3 | INT, INTEGER | 4 | Value is a signed, big-endian, 32-bit twos-complement integer. |
| 0x04 | 4 | BIGINT, LONG | 8 | Value is a signed, big-endian, 64-bit twos-complement integer. |
| 0x05 | 5 | FLOAT | 4 | Value is a single precision, big-endian, IEEE 754-2008 32-bit floating point number. |
| 0x06 | 6 | DOUBLE, REAL | 8 | Value is a double precision, big-endian, IEEE 754-2008 64-bit floating point number. |
| 0x08 | 8 | YEAR | 1 | Value is an 8-bit twos-complement integer. Both positive and negative numbers are supported in the range -128 to 127. This indicates a year with respect to the year 2000, i.e. 1872-2127 |
| 0x09 | 9 | TIME | 4 | Value is a big-endian 32-bit twos-complement integer. Indicates time of day in milliseconds since midnight, i.e. "millis". Note that only values of [0-86,400,000) are valid. i.e. Hex [0x00000000-0x005c00) |
| 0x0A | 10 | DATETIME | 8 | A big-endian unsigned LONG integer that represents the specified number of milliseconds since the standard base time known as "the epoch". It should display as a formatted string: YYYY-MM-DD hh:mm:ss, e.g. 2016-03-23 13:52:23. |
| 0x0B | 11 | DATE | 8 | A datetime whose time component is 00:00:00, but does not display. |
| 0x0C + n | 12 + n | TEXT | | Value is an ASCII string of length n. C-style string null terminators are not used or needed. |

Figure 4 - Data Types, Their 1-byte Code, and Implementation

Example



Similarity to Genes ?

